

Open Geospatial Consortium

Submission Date: <yyyy-mm-dd>

Approval Date: <yyyy-mm-dd>

Publication Date: <2020-05-27>

External identifier of this OGC® document: <http://www.opengis.net/doc/{doc-type}/{standard}/{m.n}>

Internal reference number of this OGC® document: 20-046

Category: OGC® White Paper

Editor: <Name(s) of Editor or Editors>

OGC Publish-Subscribe White Paper

Copyright notice

Copyright © <year> Open Geospatial Consortium

To obtain additional rights of use, visit <http://www.opengeospatial.org/legal/>

Warning

This document is not an OGC Standard. This document is an OGC White Paper and is therefore not an official position of the OGC membership. It is distributed for review and comment. It is subject to change without notice and may not be referred to as an OGC Standard. Further, an OGC White Paper should not be referenced as required or mandatory technology in procurements.

Document type: OGC® White Paper

Document subtype:

Document stage: Draft

Document language: English

License Agreement

Permission is hereby granted by the Open Geospatial Consortium, ("Licensor"), free of charge and subject to the terms set forth below, to any person obtaining a copy of this Intellectual Property and any associated documentation, to deal in the Intellectual Property without restriction (except as set forth below), including without limitation the rights to implement, use, copy, modify, merge, publish, distribute, and/or sublicense copies of the Intellectual Property, and to permit persons to whom the Intellectual Property is furnished to do so, provided that all copyright notices on the intellectual property are retained intact and that each person to whom the Intellectual Property is furnished agrees to the terms of this Agreement.

If you modify the Intellectual Property, all copies of the modified Intellectual Property must include, in addition to the above copyright notice, a notice that the Intellectual Property includes modifications that have not been approved or adopted by LICENSOR.

THIS LICENSE IS A COPYRIGHT LICENSE ONLY, AND DOES NOT CONVEY ANY RIGHTS UNDER ANY PATENTS THAT MAY BE IN FORCE ANYWHERE IN THE WORLD.

THE INTELLECTUAL PROPERTY IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, AND NONINFRINGEMENT OF THIRD PARTY RIGHTS. THE COPYRIGHT HOLDER OR HOLDERS INCLUDED IN THIS NOTICE DO NOT WARRANT THAT THE FUNCTIONS CONTAINED IN THE INTELLECTUAL PROPERTY WILL MEET YOUR REQUIREMENTS OR THAT THE OPERATION OF THE INTELLECTUAL PROPERTY WILL BE UNINTERRUPTED OR ERROR FREE. ANY USE OF THE INTELLECTUAL PROPERTY SHALL BE MADE ENTIRELY AT THE USER'S OWN RISK. IN NO EVENT SHALL THE COPYRIGHT HOLDER OR ANY CONTRIBUTOR OF INTELLECTUAL PROPERTY RIGHTS TO THE INTELLECTUAL PROPERTY BE LIABLE FOR ANY CLAIM, OR ANY DIRECT, SPECIAL, INDIRECT OR CONSEQUENTIAL DAMAGES, OR ANY DAMAGES WHATSOEVER RESULTING FROM ANY ALLEGED INFRINGEMENT OR ANY LOSS OF USE, DATA OR PROFITS, WHETHER IN AN ACTION OF CONTRACT, NEGLIGENCE OR UNDER ANY OTHER LEGAL THEORY, ARISING OUT OF OR IN CONNECTION WITH THE IMPLEMENTATION, USE, COMMERCIALIZATION OR PERFORMANCE OF THIS INTELLECTUAL PROPERTY.

This license is effective until terminated. You may terminate it at any time by destroying the Intellectual Property together with all copies in any form. The license will also terminate if you fail to comply with any term or condition of this Agreement. Except as provided in the following sentence, no such termination of this license shall require the termination of any third party end-user sublicense to the Intellectual Property which is in force as of the date of notice of such termination. In addition, should the Intellectual Property, or the operation of the Intellectual Property, infringe, or in LICENSOR's sole opinion be likely to infringe, any patent, copyright, trademark or other right of a third party, you agree that LICENSOR, in its sole discretion, may terminate this license without any compensation or liability to you, your licensees or any other party. You agree upon termination of any kind to destroy or cause to be destroyed the Intellectual Property together with all copies in any form, whether held by you or by any third party.

Except as contained in this notice, the name of LICENSOR or of any other holder of a copyright in all or part of the Intellectual Property shall not be used in advertising or otherwise to promote the sale, use or other dealings in this Intellectual Property without prior written authorization of LICENSOR or such copyright holder. LICENSOR is and shall at all times be the sole entity that may authorize you or any third party to use certification marks, trademarks or other special designations to indicate compliance with any LICENSOR standards or specifications. This Agreement is governed by the laws of the Commonwealth of Massachusetts. The application to this Agreement of the United Nations Convention on Contracts for the International Sale of Goods is hereby expressly excluded. In the event any provision of this Agreement shall be deemed unenforceable, void or invalid, such provision shall be modified so as to make it valid and enforceable, and as so modified the entire Agreement shall remain in full force and effect. No decision, action or inaction by LICENSOR shall be construed to be a waiver of any rights or remedies available to it.

Table of Contents

| | |
|---|----|
| 1. References | 6 |
| 2. Terms and definitions | 7 |
| 2.1. Broker Brokering Publisher | 7 |
| 2.2. Message | 7 |
| 2.3. Publication | 7 |
| 2.4. Publisher | 7 |
| 2.5. Receiver | 7 |
| 2.6. Reliable Publisher | 7 |
| 2.7. Sender | 7 |
| 2.8. Subscriber | 8 |
| 2.9. Subscription | 8 |
| 3. Conventions | 9 |
| 3.1. Abbreviated terms | 9 |
| 3.2. UML notation | 9 |
| 4. Overview | 10 |
| 5. Asynchronous Operations | 11 |
| 5.1. Behavior Models | 11 |
| 5.1.1. Request-Response | 11 |
| 5.1.2. Delayed Response | 11 |
| 5.1.3. Standing Request | 12 |
| 5.1.4. Synchronization | 13 |
| 5.1.5. Publish-Subscribe | 13 |
| 5.1.6. Broadcast | 14 |
| 5.2. Notification and Alert | 14 |
| 5.2.1. Callback | 15 |
| 5.2.2. Polling | 15 |
| 5.2.3. Stored response queue | 15 |
| 5.2.4. Man in the Loop | 15 |
| 5.3. Filtering | 15 |
| 5.3.1. Event (RSS, SNMP) | 15 |
| 5.3.2. Topic Hierarchy | 15 |
| 5.3.3. Query expression (Standing Query) | 15 |
| 5.3.4. Check Point | 16 |
| 6. OGC's PubSub Implementation Standard | 17 |
| 6.1. Overview | 17 |
| 6.2. OGC's PubSub features | 17 |
| 6.3. Related Work | 20 |
| 6.4. Basic PubSub 1.0 extension for the generic OWS | 21 |

| | |
|---|----|
| 6.4.1. Conceptual model | 21 |
| 6.4.2. Required Capabilities components | 22 |
| 6.5. Support to legacy components | 26 |
| 7. PubSub in OGC SensorThings | 30 |
| 8. W3C Pub Sub Recommendation | 32 |
| 9. Asynchronous Messaging for Aviation from OGC testing | 33 |
| 10. WMO OpenWIS use of AMQP | 34 |
| 11. AsyncAPI description | 37 |
| 12. Ideas for AsyncAPI in OGC | 38 |
| 13. Ideas for an OGC Abstract Spec for PubSub | 39 |
| Annex A: Title | 40 |
| Annex B: Revision History | 41 |
| Annex C: Bibliography | 42 |

i. Abstract

<Insert Abstract Text here>

ii. Keywords

The following are keywords to be used by search engines and document catalogues.

ogcdoc, OGC document, <tags separated by commas>

iii. Preface

NOTE

Insert Preface Text here. Give OGC specific commentary: describe the technical content, reason for document, history of the document and precursors, and plans for future work. > Attention is drawn to the possibility that some of the elements of this document may be the subject of patent rights. The Open Geospatial Consortium shall not be held responsible for identifying any or all such patent rights.

Recipients of this document are requested to submit, with their comments, notification of any relevant patent claims or other intellectual property rights of which they may be aware that might be infringed by any implementation of the standard set forth in this document, and to provide supporting documentation.

iv. Submitting organizations

The following organizations submitted this Document to the Open Geospatial Consortium (OGC):

Organization name(s)

v. Submitters

All questions regarding this submission should be directed to the editor or the submitters:

Name Affiliation

Chapter 1. References

The following documents are referenced in this document. For dated references, subsequent amendments to, or revisions of, any of these publications do not apply. For undated references, the latest edition of the normative document referred to applies.

- OGC 06-121r3, OGC® Web Services Common Specification, version 1.1.0 (9 February 2007)
- OGC 13-131, OGC® Publish/Subscribe Interface Standard 1.0 - Core
- OGC 13-133, OGC® Publish/Subscribe Interface Standard 1.0 - SOAP Protocol Binding Extension
- OGC 13-132, OGC® Publish/Subscribe Interface Standard 1.0 - RESTful Protocol Binding Extension (draft, in progress)
- OGC 07-006r1, OpenGIS® Catalogue Services Specification, version 2.0.2, corrigendum 2
- OGC 07-045, OpenGIS® Catalogue Services Specification 2.0.2 - ISO Metadata Application Profile, version 1.0
- OGC 07-110r4, CSW-ebRIM Registry Service - Part 1: ebRIM profile of CSW, version 1.0.1, corrigendum 1
- OGC 07-144r4, CSW-ebRIM Registry Service - Part 2: Basic extension package, version 1.0.1, corrigendum 1
- OGC 08-103r2, CSW-ebRIM Registry Service - Part 3: Abstract Test Suite, version 1.0.1
- OGC 16-137r2, Testbed-12 PubSub / Catalog Engineering Report (12 May 2017)

Chapter 2. Terms and definitions

For the purposes of this report, the definitions specified in Clause 4 of the OWS Common Implementation Standard [OGC 06-121r3] shall apply. In addition, the following terms and definitions apply.

2.1. Broker | Brokering Publisher

Intermediary between Subscribers and other Publishers which have been previously registered with the broker. The broker is not the original producer of messages, but only acts as a message middleman, re-publishing messages received from other Publishers and decoupling them from their Subscribers

2.2. Message

A container within which data (such as XML, binary data, or other content) is transported. Messages may include additional information beyond data, including headers or other information used for routing or security purposes

2.3. Publication

A uniquely identified aggregation of messages published by a Publisher over time. A Publisher may offer any number of publications that Subscribers may subscribe to

2.4. Publisher

An entity that offers publications to Subscribers; supports subscription management (subscribe, unsubscribe) and is responsible for filtering and matching messages of interest to active subscriptions

2.5. Receiver

An entity that receives messages from Senders; may (but need not) be the original Subscriber

2.6. Reliable Publisher

A publisher of messages that offers capabilities to detect and recover from message delivery losses, whether caused by network failures, software failures, hardware failures, or other causes

2.7. Sender

Entity that sends messages to Receivers; may (but need not) be the initial creator/producer of the data in the message payload

2.8. Subscriber

Entity that creates a subscription at a Publisher; may (but need not) be the Receiver of delivered messages

2.9. Subscription

Expression of interest in all or part of a publication offered by a Publisher. When a subscription has been created, the Publisher delivers messages that match the subscription criteria to the Receiver defined in the subscription

Chapter 3. Conventions

3.1. Abbreviated terms

- CFP Call for Participation
- DWG Domain Working Group
- HTTP Hypertext Transfer Protocol
- MEP Message Exchange Pattern
- OAI-PMH Open Archives Initiative Protocol for Metadata Harvesting
- OGC Open Geospatial Consortium
- RFQ Request for Quotation
- SSE Server-Sent Events
- SWG Standards Working Group
- UML Unified Modeling Language
- XML eXtensible Markup Language

3.2. UML notation

Most diagrams that appear in this document are presented using the UML 2 static structure diagram, as described in Subclause 5.2 of [OGC 06-121r9].

All classes in this document are extensible and may be extended with application- or domain-specific content via Extension blocks.

NOTE

The UML shown in this document is considered conceptual and abstract, and should not be interpreted as an implementation strategy for bindings that extend and implement a standard. For example, `TM_Instant` from ISO 19108 may be used to represent time instants for conceptual clarity, but bindings and implementations of this document may realize `TM_Instant` as a GML `TimeInstant`, an ISO 8601 date string, or any other representation that is consistent with `TM_Instant`.

Chapter 4. Overview

Chapter 5. Asynchronous Operations

An asynchronous operation is any interaction between two software entities where the concluding action does not immediately follow the initiating action. There are a number of different models for Asynchronous Operations. This section attempts to describe some of the more common ones. An understanding of the scope of this space will help discuss and compare the alternatives OGC faces in asynchronous services.

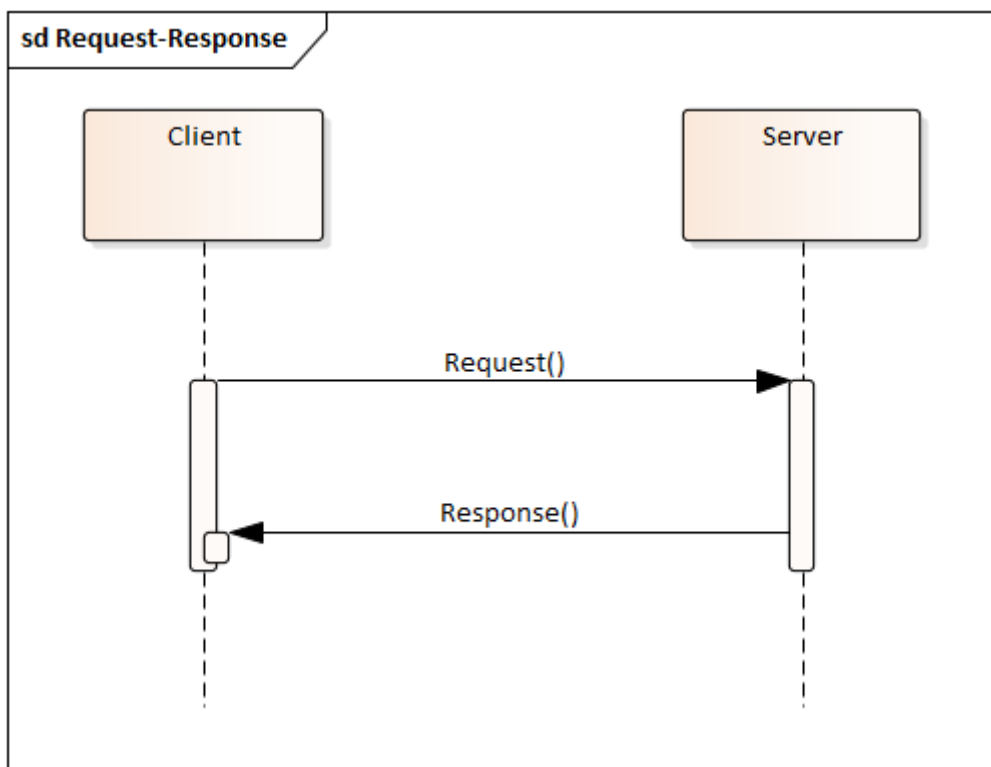
This section will look at three aspects of asynchronous operations:

1. The overall pattern of behavior
2. How the two entities rendezvous
3. The criteria for selecting response messages

5.1. Behavior Models

5.1.1. Request-Response

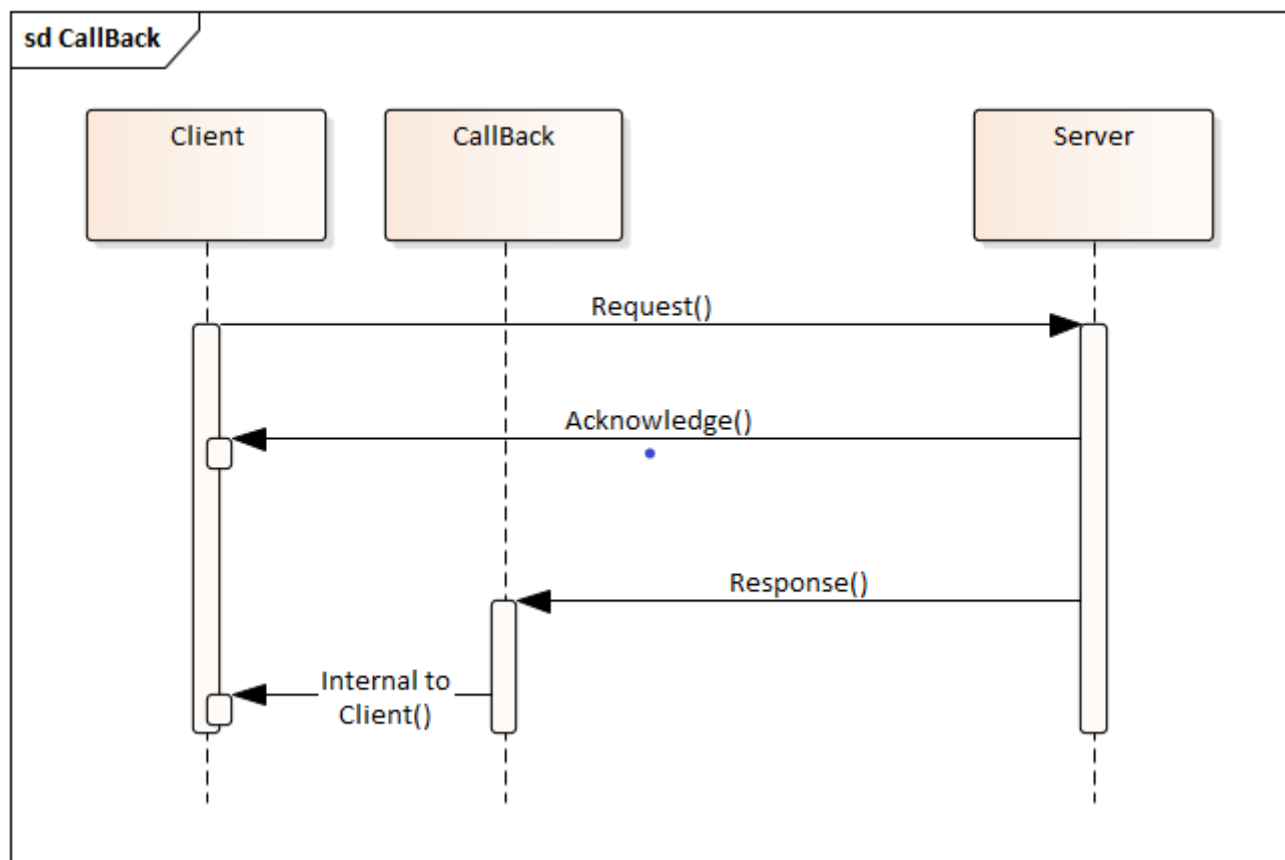
Request-Response is a synchronous behavior model. A request is issued by one entity and a response is provided in return. Asynchronous behaviors can best be understood in contrast to this model.



5.1.2. Delayed Response

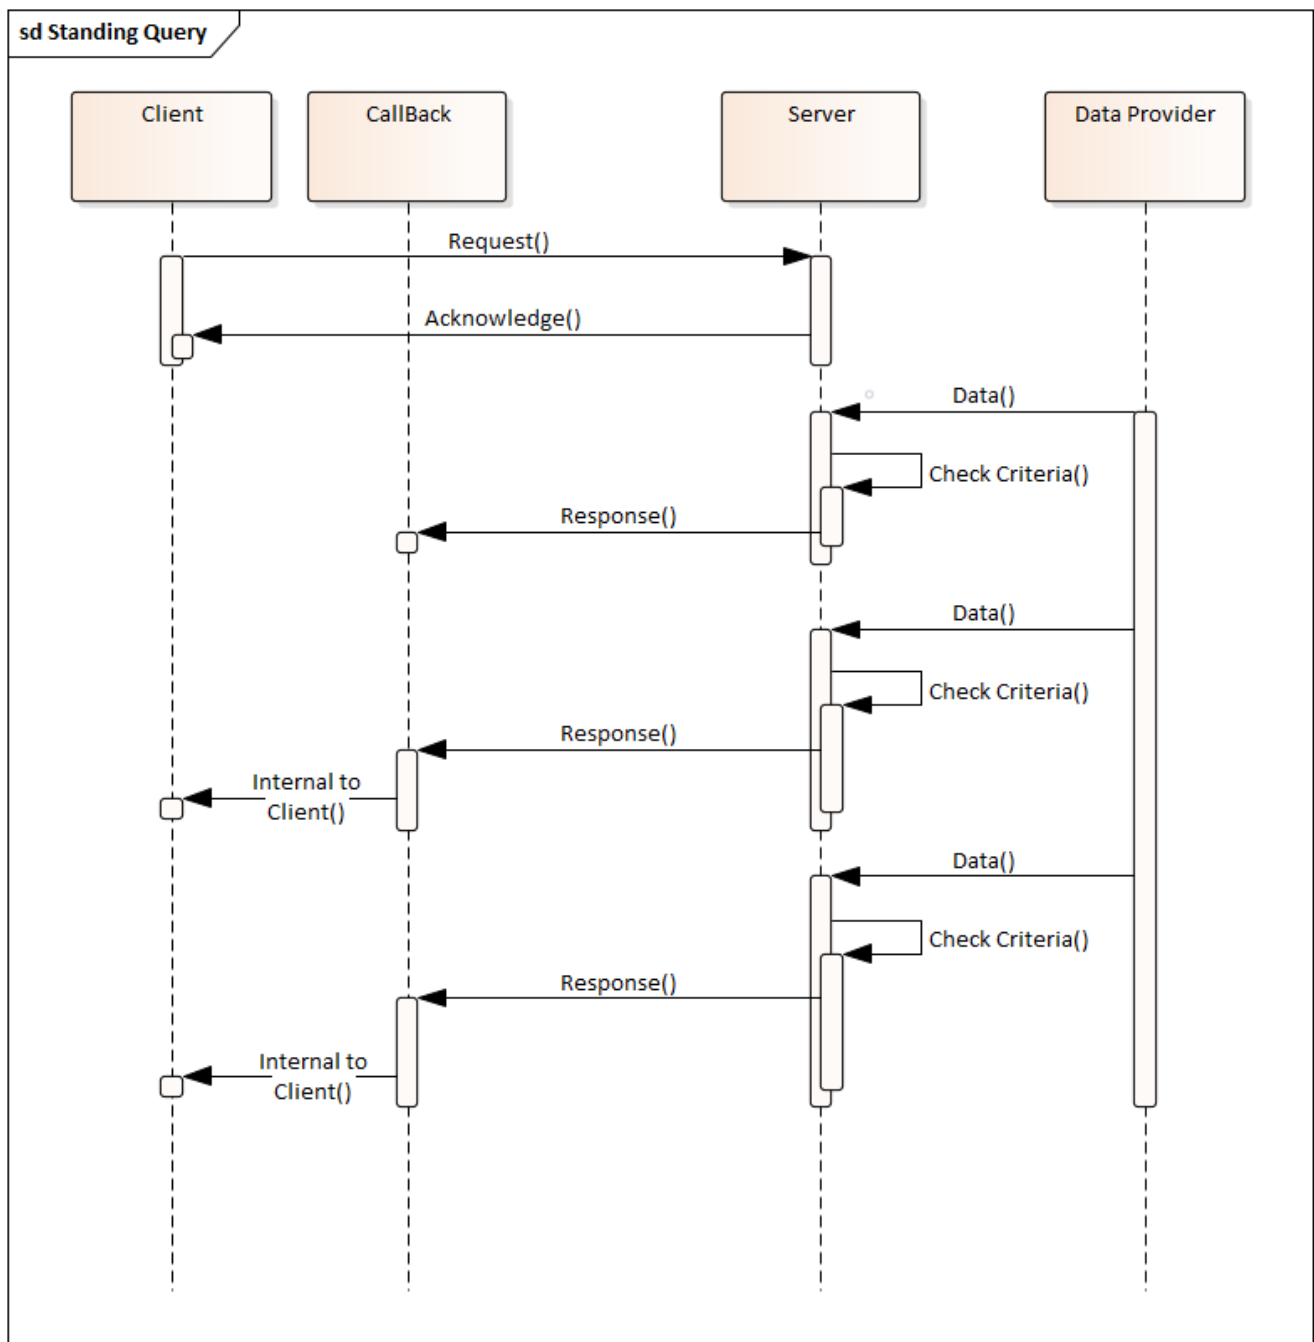
This model differs from Request-Response in that the immediate response is not the final response. Rather it acknowledges successful receipt of the request. The final response is to be delivered later.

A key part of this pattern is the callback. A callback is an http accessible resource which can receive information (responses) from the server.



5.1.3. Standing Request

A Standing Request is a variation of the Delayed Response pattern. While a Delayed Response performs a single operation, a Standing Request is active until instructed to stop.



5.1.4. Synchronization

The Synchronization pattern supports a scenario where communication between the message producer and consumer is intermittent. When communication is possible, they perform whatever transactions are needed to synchronize their states, then establish a checkpoint for that state. Both parties can then continue to operate independently until the next synchronization opportunity arrives.

5.1.5. Publish-Subscribe

A messaging pattern where senders of messages, called publishers, do not program the messages to be sent directly to specific receivers, called subscribers, but instead to categorize published messages into classes without knowledge of which subscribers, if any, there may be. Similarly, subscribers express interest in one or classes and only receive messages that are of interest,

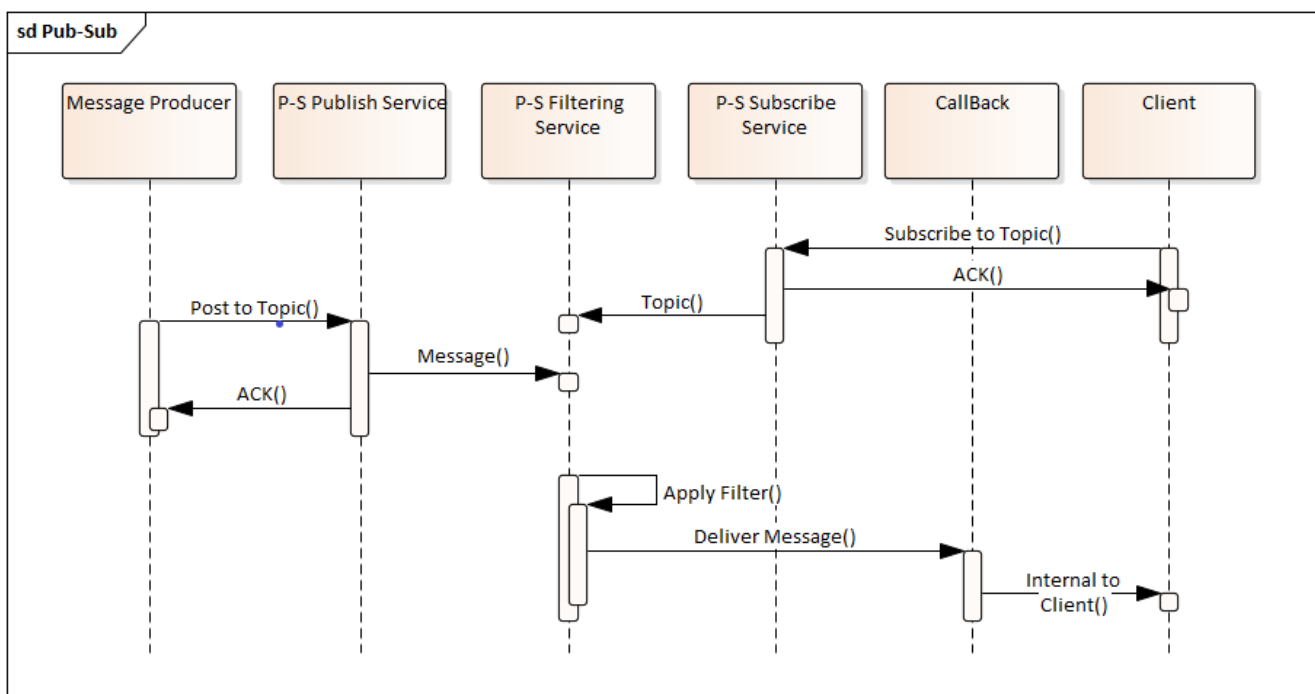
without knowledge of which publishers, if any, there are. (Wikipedia)

The Publish/Subscribe model is distinguished from the request/reply and client/server models by the asynchronous delivery of messages and the ability for a Subscriber to specify an ongoing (persistent) expression of interest.

Alternate Text

The Publish-Subscribe model completely separates message producers and consumers. Potential consumers of messages create filtering criteria which describe the types of messages they wish to receive. They then "subscribe" to a Pub-Sub service with this filtering criteria. Producers of messages "publish" those messages to the Pub-Sub service along with a set of tags which describe each message. The Pub-Sub service evaluates the tags against the filtering criteria of all subscribers. The message is forwarded to all subscribers who's criteria are met.

The "publish" operation follows the Request-Response pattern. The "subscribe" operation follows the Standing Request pattern.



5.1.6. Broadcast

Broadcast is the simplest asynchronous pattern. The message producer simply sends the message to everyone. It is left up to the recipients to decide what to do with it.

5.2. Notification and Alert

An inherent property of Asynchronous operations is that there is no persistent connection between the message producer and message receiver. Therefore, there must be a way for the message producer to re-establish a connection with the receiver in order to complete the transaction. There are a number of ways this is done.

5.2.1. Callback

Callbacks can be viewed as mini-services who's sole purpose is to receive an asynchronous response. Information on how to access the callback is provided with the initial request. Message producers (or their agents) use this information deliver responses, typically using the Request-Response pattern.

5.2.2. Polling

In polling the requesting entity checks on the status of their request on a recurring basis. Upon completion of the request, the requestor retrieves the result to complete the transaction.

5.2.3. Stored response queue

A stored response queue is a service which holds responses to asynchoronous requests. The message producer simply leaves the response in the queue, and it's up to the requestor to retrieve it.

5.2.4. Man in the Loop

If all else fails, let the human do it. Many alternatives are available including instand messaging, e-mail, phone calls, even the Postal Service.

5.3. Filtering

Filtering allows a message producer to identify the intended recipients of a message.

5.3.1. Event (RSS, SNMP)

Event filtering specifies that a notification will be sent if certain conditions occure. For example, if the free space in a mail box drops below 10%.

5.3.2. Topic Hierarchy

Publish-Subscribe implementations typically define a set of topics (terms) which can be used to select messages for delivery. In the most basic case a recipient can only subscribe to topics. More capable systems may provide a simple query language to go with the topic vocabulary.

Example: MQTT uses Topic Filters to select messages. A Topic Filter is a path-like hierarchy of concepts. Wildcards are supported to indicate a single path entry or multiple. For example:

1. sport/tennis/player1/score/Wimbledon is a Topic Name
2. sport/+/player1 is a Topic Name with a wild card for only one level
3. sport/tennis/#/ranking is a Topic Name with a wild card for 1 or more levels.

5.3.3. Query expression (Standing Query)

Java Messaging Service (JMS) is the foundation for many (most) publish-subscribe services. JMS

supports messaging selection through a query string. The query language is a subset of the SQL92

More capable systems support a full query language for filtering messages. For example, an asynchronous WFS would accept asynchronous requests using the same Filter Encoding language as any other WFS. But the results would be returned asynchronously.

5.3.4. Check Point

A check point is a store snapshot of the state of the system as a specific date and time. All changes made after a check point are can

Chapter 6. OGC's PubSub Implementation Standard

6.1. Overview

The OGC has conducted significant work on event-based models and architectures. Figure 1 provides an overview and a timeline of relevant OGC work in the last decade.

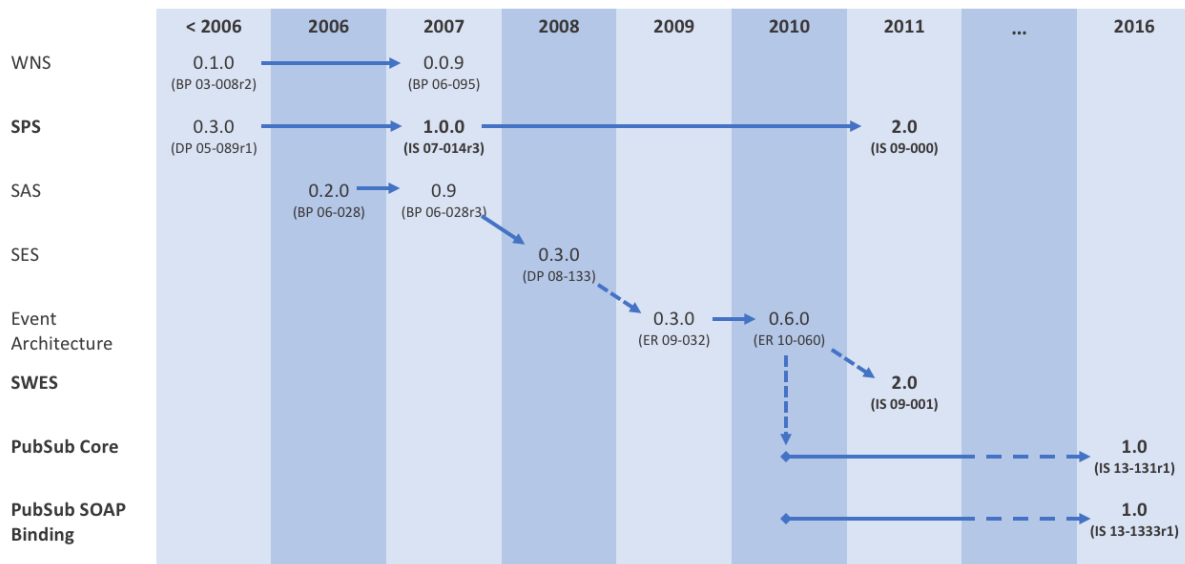


Figure 1. OGC Publish/Subscribe related work (source: [Rieke2018]). The rows show the roadmap of various documents, including official OGC standards (highlighted in bold fonts) and milestones, indicated with their version and OGC reference number and type (Best Practice, Discussion Paper, Implementation Standard, and Engineering Report; see section [Abbreviated terms](#) for a complete list of abbreviations)

From 2010 onward, these efforts have been subsumed under the scope of the PubSub SWG and resulted in the adoption of the Publish/Subscribe Interface Standard 1.0 in February 2016.

6.2. OGC's PubSub features

The OGC Publish/Subscribe Interface Standard (in short, PubSub) describes a mechanism to support publish/subscribe requirements across OGC service interfaces, such as Sensor Observation Service (SOS) and Catalogue Service for the Web (CSW), and data types, such as coverages, features, and observations.

PubSub is intended as an overarching model to complement OGC service interfaces with publish/subscribe capabilities, alongside the primarily addressed request/reply model. In fact, the latter is sufficient to meet many use cases, when clients request the data of interest upon need, and may re-issue requests periodically for updates. However, the publish/subscribe model can be better suited to manage urgent, unpredictable pieces of information, such as sporadic events, hazard warnings and sensor data, resulting in less network traffic and more timely responses.

In the assumption that clients are interested in the information already provided by OGC services, the rationale for PubSub is to avoid impact on the existing service semantics, when possible reusing the specialized mechanisms for accessing and filtering the information provided (e.g., the Filter

Encoding Specification ^[1] for general-purpose filtering). Hence, PubSub provides notification and messaging capabilities as a module to existing services, as well as a standalone, independent OWS capability. Besides, it supports simple notifications (e.g. references to data from OWS services) or actual data push.

The basic OGC PubSub workflow, depicted in Figure 2, involves two primary actors: the Publisher of information (XML, binary data, or other content); and the Subscriber, which subscribes to the Publisher to obtain the published information, possibly filtered according to criteria defined upon subscription. Information is contained and transported in messages, which may include additional data, such as for routing or security purposes.

PubSub distinguishes two additional roles, the Sender and the Receiver, respectively the entity that delivers messages and the entity to which messages are delivered, allowing versatile use cases (e.g. a Subscriber entity may subscribe to a Publisher a number of third parties acting as Receivers; or a Publisher may be segregated from the ultimate recipients of its messages by a separate Sender).

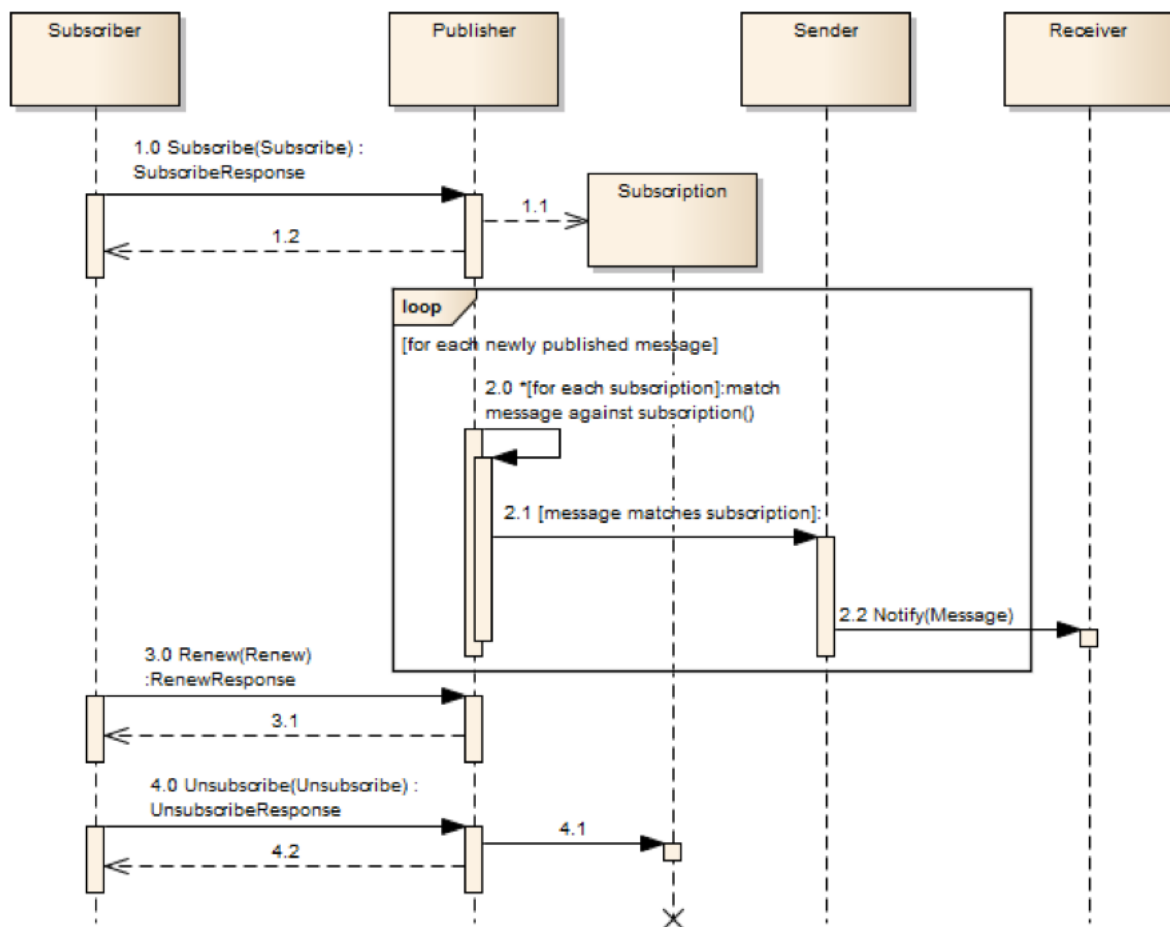


Figure 2. OGC PubSub basic workflow (source: [PubSubCore2016]). The Subscriber creates a subscription (with possible filter criteria) on behalf of a Receiver using the Subscribe operation on a Publisher (1.0). The Publisher creates a Subscription (1.1) and returns a SubscribeResponse (1.2). The Publisher attempts to match each new message against the filter of each Subscription (2.0). If a message matches, the Sender delivers it to the appropriate Receiver via the Notify operation (2.1). The Subscriber may utilize the Renew operation (3.0) to extend the lifetime of a Subscription. If the Publisher accepts the request, it sets the new termination time on the Subscription and returns a RenewResponse (3.1). The Subscriber may at any time request the termination of a Subscription via the Unsubscribe operation (4.0). If the Publisher accepts the request, it terminates the subscription (4.1) and returns an UnsubscribeResponse (4.2).

To qualify as a Publisher, an OWS must at least conform to the Basic Publisher conformance class, that is satisfy the requirements for the above basic functionalities. Besides, it must return in its GetCapabilities response the three additional Capabilities components represented in [Figure 3](#).

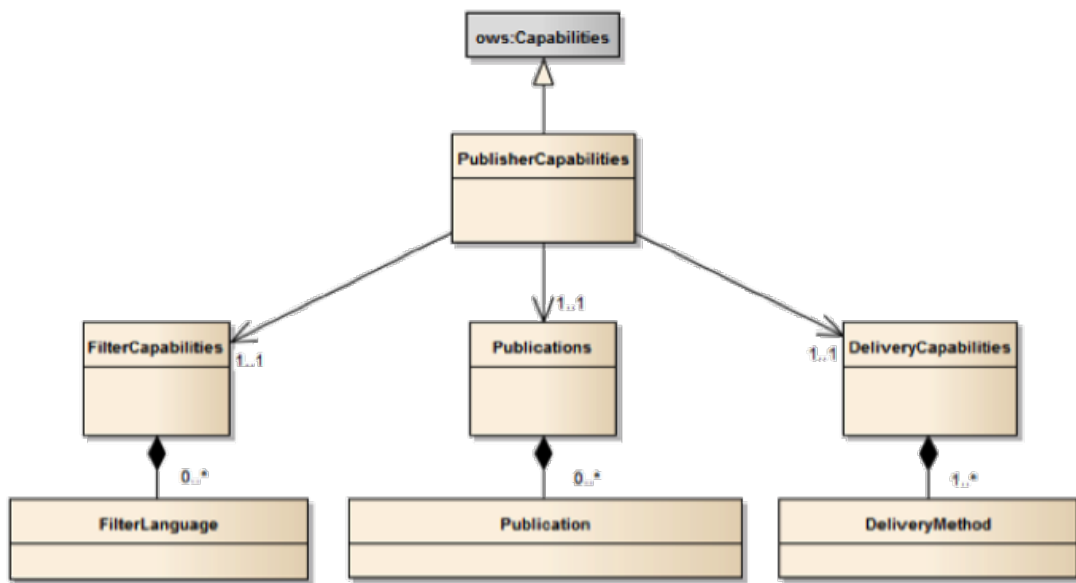


Figure 3. OGC PubSub additional Capabilities components (source: [\[PubSubCore2016\]](#)).

The Publications component describes the contents offered by the Publisher, i.e. the sets of messages that Subscribers can subscribe to. The specification is agnostic as to what constitutes a publication, i.e. what events should cause notifications by a Publisher (i.e. its event model). Likewise, PubSub is agnostic as for the encoding of messages.

The FilterCapabilities component describes the filtering-related capabilities of a Publisher, i.e. the filter languages it supports for matching messages against subscriptions (e.g., OGC Filter Encoding, XQuery). PubSub is agnostic as for the language to filter messages in subscriptions. A Publisher may support multiple filter languages, to support different Subscribers.

The DeliveryCapabilities component describes the methods supported by the PubSub-enabled OWS for delivering messages. The publish/subscribe MEP typically implies push-style message delivery, however some delivery methods may actually be underpinned by pull-based mechanisms (e.g. polling). Examples of delivery methods include: SOAP and related technologies, such as WS-N (used by the PubSub SOAP Binding), ATOM, PubSubHubbub, OAI-PMH, e-mail, Short Message Service, WebSockets and SSE. The PubSub standard is agnostic as regards delivery methods. A Publisher may offer more than one method of delivery for each Publication, to be chosen by Subscribers.

In addition to the mandatory Basic Publisher conformance class, PubSub defines several other optional conformance classes, introducing additional functionalities, e.g. to pause a Subscription (Pausable Publisher), derive additional publications (Publication Manager), group messages in batches (Message Batching Publisher). Moreover, it defines conformance classes to bind such functionalities to actual technologies. At present, the PubSub specification consists of two parts:

1. a Core document [\[PubSubCore2016\]](#) that abstractly describes the basic mandatory functionalities and several optional extensions, independently of the underlying binding technology;
2. a SOAP binding document [\[PubSubSOAP2016\]](#) that defines the implementation of PubSub in

SOAP services, based on the OASIS Web Services Notification (WS-N) set of standards. [WSBN2016]

The scope of the OGC PubSub Standard Working Group also includes a RESTful binding, to realize the PubSub functionality in REST/JSON services. Several communities are proposing additional extensions (e.g. bindings to JMS and MQTT), leveraging on the modular structure of the OGC specifications.

PubSub Core requires that a PubSub-enabled OWS advertise the implemented Conformance Classes in its Capabilities document. Figure 4 shows all the conformance classes currently defined by the PubSub specification.

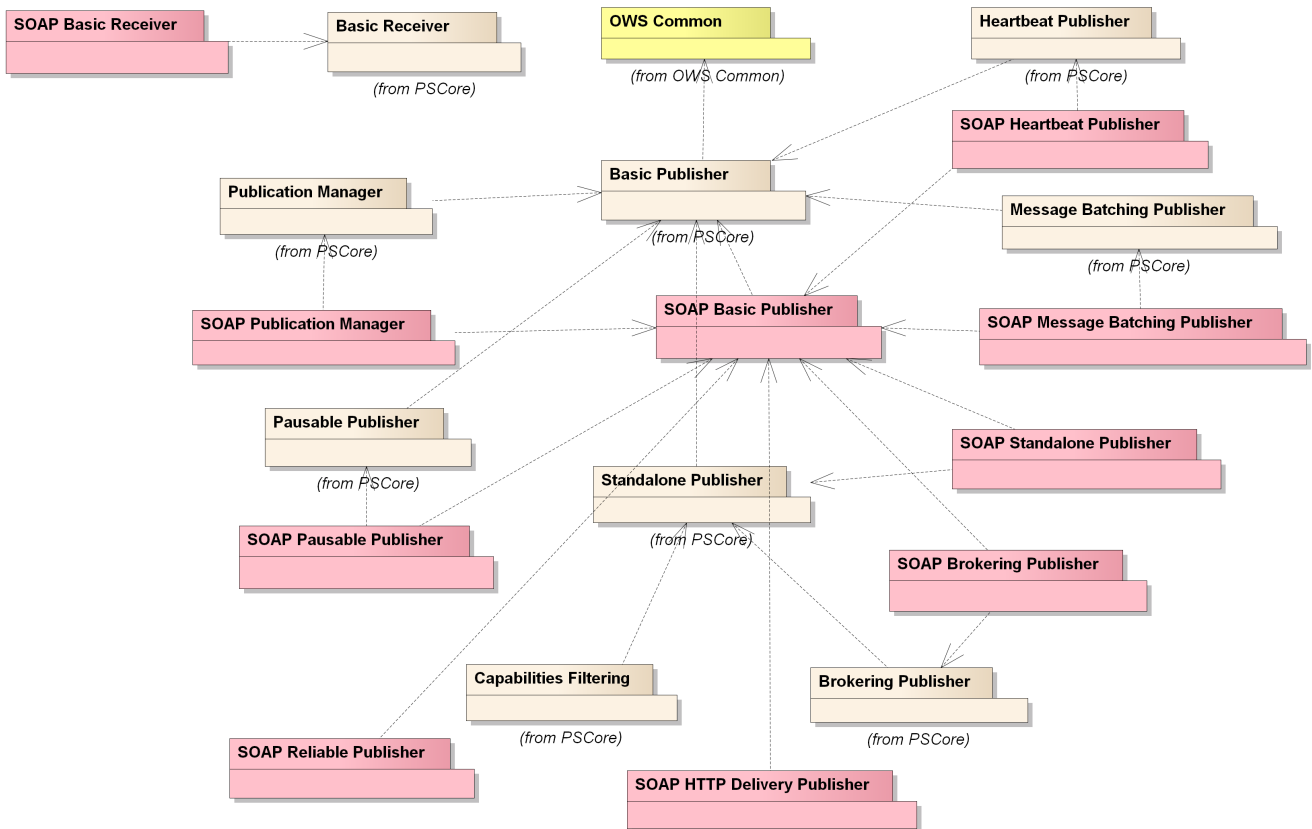


Figure 4. OGC PubSub Conformance classes (source: [PubSubCore2016]).

6.3. Related Work

Recognizing that the OGC baseline mainly supported synchronous web service request-response capabilities, the 2016 OGC Testbed 12 initiative addressed the means to incorporate forms of asynchronous service interaction, including Publish/Subscribe message patterns, for example in WPS, WCS, WFS, or in application domains such as the Sensor Web and Aviation. [Bigagli2017]

In particular, the RFQ/CFP [2] included a specific Asynchronous Service Interaction subtask, part of a set of subtasks that aimed at enhancing the OGC Baseline, by extending OGC architectural designs through efforts that cross over several individual standards and services and are applied in a much wider scope.

The subtask description in the RFQ/CFP distinguished among three different approaches to handle asynchronous interaction with OGC Web services:

1. WPS façades;
2. Specific extensions to each OGC Web Service with asynchronous request/response capabilities;
3. OGC PubSub.

The document deliverable "A067 Implementing Asynchronous Service Response Engineering Report" (OGC 16-023) elaborates on the above approaches in situations where big chunks of data require asynchronous delivery. The ER focuses on the first and the second approach, with the goal to summarize and compare the results from using a WPS facade and an extension for WFS for asynchronous service responses, as well as to provide recommendations for future activities.

The document deliverable "A074 PubSub/Catalog Engineering Report" (OGC 16-137) [\[OGC19-137r2\]](#) focused on the third approach, OGC PubSub, and exemplified the use of the standard, particularly in conjunction with the Catalog Service interface, investigating the functional requirements of an interoperable, push-based data discovery solution. As underlined in the RFQ/CFP, it is important to provide methods that support notification (push) of new data as opposed to search (pull), given the volume of data typically available in catalogs.

Besides, it introduced a general, basic mechanism for enabling PubSub for the generic OGC Web Service over the existing request/reply OWS's, i.e. usual requests as filters, usual responses as appropriate updates/data pushes, existing semantics and syntax expressiveness. The following chapters summarize such mechanism.

6.4. Basic PubSub 1.0 extension for the generic OWS

The PubSub extension for the generic OWS introduced by OGC 16-137 is conceived as a simple way to enable the existing request/reply OWS specifications to Publish/Subscribe, by implementing the OGC Publish/Subscribe Interface Standard 1.0.

An OWS implementing this extension is capable of accepting its usual requests as filters, and of sending notifications about data/metadata updates, based on its existing semantics and syntax expressiveness.

6.4.1. Conceptual model

This chapter describes how PubSub 1.0 Core operations, encodings and messages are modeled in terms of the functionalities of the generic OWS. No assumption is made on the capabilities of the target OWS, other than those defined by the OGC Web Services Common Standard. Hence this extension may apply, for example, to WFS, WCS, and other OWS interfaces.

The PubSub specification is agnostic as to what constitutes a change, i.e. an event that should cause a notification by a Publisher (aka its event model). It is only required that a Publisher instance communicate what notifications it will emit by advertising them in the Publication section of its Capabilities document (see below).

In general, a PubSub-OWS may be able to notify about changes to any component of its information set. For example, it may notify about changes to its Capabilities document. The extension introduced in this chapter addresses the most general case, at the expenses of efficiency and semantic accuracy. The precise definition of an event model for the various OWS's is left to the

relevant OGC Working Groups.

The basic PubSub-OWS MEP can be generalized as follows (see [Figure 5](#)):

1. The OWS client subscribes specifying a request to be used as filter for the notifications;
2. The OWS client obtains the Time-0 response via a standard Request/Reply, with the same request as above;
3. The OWS notifies the client of subsequent updates to the response, according to its existing semantics and syntax.

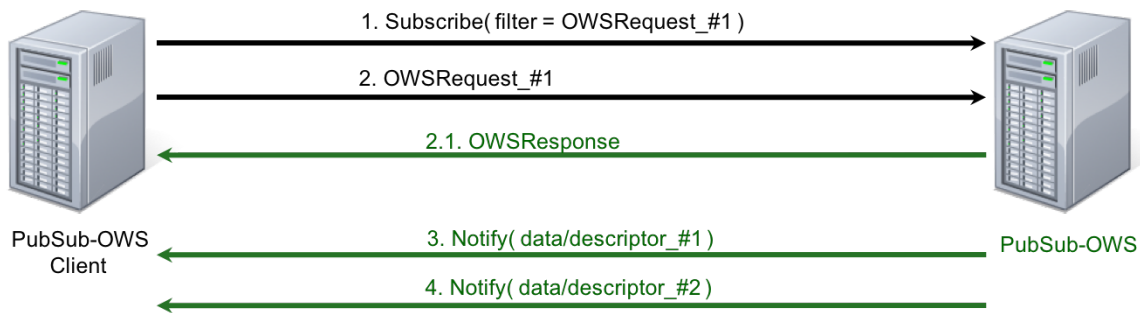


Figure 5. OWS Publish/Subscribe MEP

This may be formalized in an “OWS Request/Reply Publisher” Conformance Class that:

- Accepts OWS requests as subscription filters
 - The Publisher may constraint the filter expressions allowed in Subscriptions (e.g. by imposing OpenSearch templates)
- Sends corresponding OWS responses to notify about data/metadata updates

This MEP is a simple way to enable existing OWSs to PubSub, allowing to bind the PubSub 1.0 Core operations, encodings and messages to the standard OWS functionalities, data models, and semantics.

6.4.2. Required Capabilities components

PubSub Core requires that the OWS advertise the implemented Conformance Classes in its Capabilities document, namely in the Profile property of the ServiceIdentification section (as of OWS Common 1.1). Besides, it requires that the additional Capabilities components represented in [Figure 6](#) are returned in the GetCapabilities response, but does not specify the specific mechanism for incorporating these additional Capabilities components into the OWS Capabilities document. These extension proposes to include these additional Capabilities components in the ExtendedCapabilities of the OWS, as detailed in the following chapters.

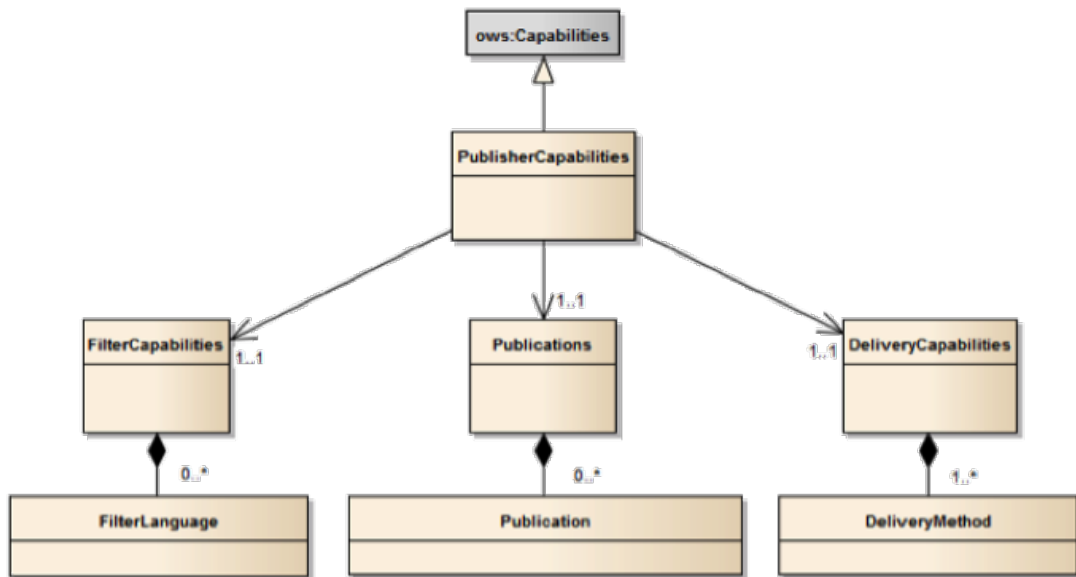


Figure 6. PubSub Capabilities components

FilterCapabilities

The FilterCapabilities section describes the filtering-related capabilities of a PubSub-OWS, i.e. the filter languages it supports for matching events against subscriptions (e.g., OGC Filter Encoding). This allows the pluggability of filter languages.

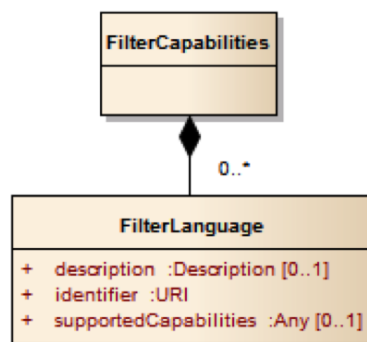


Figure 7. Filter Capabilities

The SupportedCapabilities elements allows restricting the acceptable requests, possibly providing templates. The following Capabilities snippet declares that this PubSub-OWS instance (namely, a CSW) accepts as subscription filters GetRecords requests conforming to the specified OpenSearch template. Multiple templates may be introduced, specifying multiple FilterLanguages.


```

<FilterCapabilities>
  <FilterLanguage>
    <Abstract>This PubSub-OWS accepts requests as subscription filters, according to
the OpenSearch template specified in SupportedCapabilities.
    </Abstract>
    <Identifier>http://www.opengis.net/spec/pubsub/1.0/conf/ows/request-reply-
publisher</Identifier>
    <SupportedCapabilities>http://tb12.essi-lab.eu/pubsub-
csw/services/opensearch?ct={count?}&st={searchTerms?}&bbox={geo:box?}&ts={
time:start?}&te={time:end?}
    </SupportedCapabilities>
  </FilterLanguage>
</FilterCapabilities>

```

DeliveryCapabilities

The DeliveryCapabilities section describes the delivery methods supported by the PubSub-OWS, e.g. SOAP, WS-Notification, ATOM, SSE, WebSockets, OAI-PMH. This allows the pluggability of delivery methods.




Figure 8. Delivery Capabilities

The following Capabilities snippet declares that this PubSub-OWS instance delivers notifications via SSE (see chapter [Delivery methods](#), below).

DeliveryCapabilities

```

<DeliveryCapabilities>
  <DeliveryMethod>
    <Abstract>This PubSub-OWS supports notification delivery via SSE.
    </Abstract>
    <Identifier>http://www.w3.org/TR/eventsourcing/
    </Identifier>
  </DeliveryMethod>
</DeliveryCapabilities>

```


Delivery methods

The `DeliveryCapabilities` section describes the methods supported by the PubSub-OWS for delivering notifications. Publishers may offer more than one method of delivery for each Publication, to be chosen by Subscribers. Publish/Subscribe would imply push-style message delivery, however some methods may actually be pull-based (e.g. polling), under the hood.

Examples include: SOAP and related technologies, such as WS-Notification (used by PSSB), ATOM (polling using the “If-Modified-Since” and “start-index” parameters), PubSubHubbub, OAI-PMH (polling using the “from” parameter), e-mail, SMS, WebSockets, SSE.

Server-Sent Events (SSE) is a pure push-style communication technology based on HTTP and the SSE EventSource API standardized as part of HTML5 by the W3C. A SSE client (e.g. all modern HTML 5.0 browsers) receives automatic updates from a server via HTTP connection, simply setting the following parameters:

- `ContentType`: "text/event-stream;charset=UTF-8"
- `Cache-Control`: "no-cache"
- `Connection`: "keep-alive"

Publications

The Publications section describes the contents offered by the PubSub-OWS, i.e. the sequences of notifications that Subscribers can subscribe to.




Figure 9. Publications

The following Capabilities snippet declares a publication that notifies on all the relevant events for this PubSub-OWS. Notifications can be filtered with the semantics of the requests of this OWS and are delivered via SSE, encoded in JSON (see chapter [Notification encoding](#), below).

```

<Publications>
  <Publication>
    <Abstract>>This publication notifies on all the relevant events for this PubSub-
    OWS.
    </Abstract>
    <Identifier>ALL</Identifier>
    <ContentType>application/json</ContentType>

    <SupportedFilterLanguage>http://www.opengis.net/spec/pubsub/1.0/conf/ows/request-
    reply-publisher</SupportedFilterLanguage>

    <SupportedDeliveryMethod>http://www.w3.org/TR/eventsource/</SupportedDeliveryMethod>
  </Publication>
</Publications>

```

Notification encoding

For the generic OWS instance, no operation is defined that provides the basic semantics of “insert”, “update”, and “delete” actions on the content managed by the instance.

The most generic mechanism to notify about updates is that the Publisher re-send the whole response element corresponding to the request used as filter in the Subscription. For example, in the case of WFS, if the client subscribes with a wfs:GetFeature request as a filter, the PubSub-WFS should notify about any changes by delivering a standard wfs:FeatureCollection, in response to that request.

By receiving the new response and comparing it with the previous one, a Subscriber can figure out the changes. Future evolutions of this extension may evaluate more efficient and semantically accurate encoding of notifications. A possible option for XML-based content types is XMLdiff (e.g. XML Patch, RFC 5261), or annotations (XML attributes) to add simple CRUD semantics on top of the existing XSDs.

6.5. Support to legacy components

The integration of legacy components in an eventing architecture is desirable in a number of scenarios. However, legacy components may not be instrumented to monitor their state for the purpose of notification, nor to react upon notifications from other components (or they may, but by legacy, non-standard mechanisms).

Implementing the PubSub 1.0 Standard in a legacy component may not be feasible or practical. In some cases, the legacy component can be adapted to the Publish/Subscribe MEP by an additional functional entity that realize the Publish/Subscribe functionalities. Such mediating entity acts as a proxy/adapter, i.e. a middleman between the source and the target of the message exchange, implementing the behavior and/or the interfaces required by the PubSub specification.

This use case has been considered in the phase of requirement analysis for the PubSub 1.0 standard ^[3] and is supported by the Brokering Publisher Conformance Class of the PubSub 1.0 Standard.

Depending on the intended role of the legacy component, the use case is twofold:

- Proxied Subscribe – a proxy/adaptor component subscribes to a Publisher on behalf of the legacy system and acts appropriately upon receiving notifications of interest.

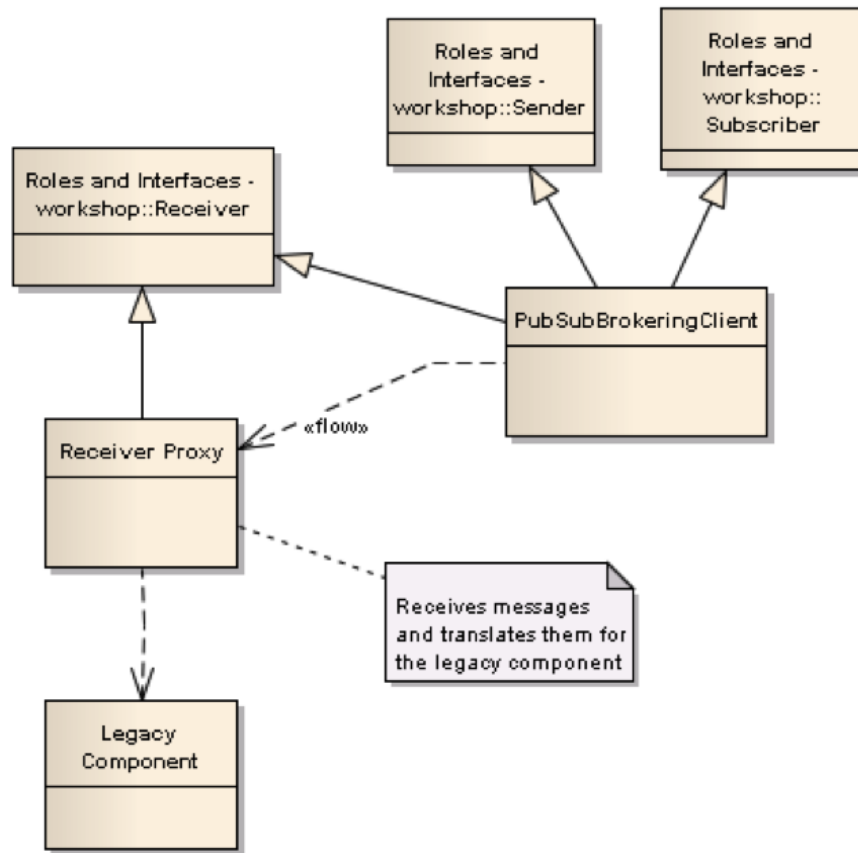


Figure 10. Proxied subscribe

- Proxied Publish – a proxy/adaptor component monitors the legacy system and generates appropriate notifications upon relevant events (according to a given event model). The proxy/adaptor may act as a full-fledged Publisher, accepting Subscriptions against the sequence of notifications, or just act as a pure Sender, relaying each notification to another Publisher entity (see [Figure 11](#)).

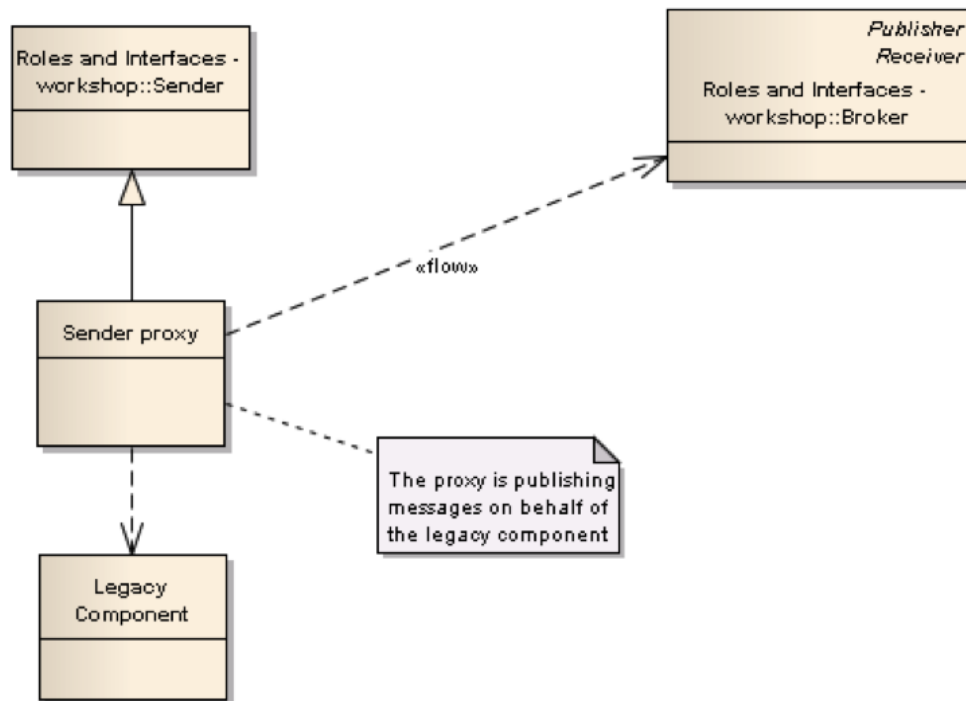


Figure 11. Proxied publish

The Brokering Publisher Conformance Class of the PubSub 1.0 Standard supports this use case. In fact, a Brokering Publisher (or, more simply, a broker), is an intermediary between Subscribers and other Publishers which have been previously registered with the broker. The broker is not the original producer of messages, but only acts as a message middleman, re-publishing messages received from other Publishers and decoupling them from their Subscribers. A broker may shuffle or aggregate messages into different publications, may offer publications with different delivery methods than the original ones, or otherwise process the messages (e.g. converting their format). A broker may also provide advanced messaging features, such as load balancing.

In general, a broker is a distinct third party that acts as a communication intermediary between the source and the target of a communication, mediating their interfaces and in some cases adding new behavior. Hence, a broker may conveniently act as a proxy/adaptor for one or more legacy components, flexibly implementing any combination of the above twofold use case.

The Brokering Publisher Conformance Class does not mandate any specific behavior to be implemented, in particular as regards the support to Delivery Capabilities, Filtering Capabilities, and Publications of the brokered Publishers. Brokers are free to interact with the brokered Publishers as appropriate for their specific application. Interactions may include subscribing to the offered publications, harvesting the data, decorating the capabilities, or other behavior (future extensions of the Conformance Class may standardize the behavior of Brokering Publishers in specific application scenarios).

Examples of Brokering Publisher applications include the following:

- Publisher Aggregation – a broker subscribes to several Publishers and relays their publications (without modifications) to interested Subscribers, acting like a Proxy to multiple Publishers. Optionally, the broker may adapt the service interface (binding) of the aggregated Publishers.
- Publication Aggregation – a broker receives messages generated by several Publishers (e.g. dumb sensors) and publishes them to the interested Subscribers as a single publication at a

single endpoint, for the sake of simpler connectivity, or improved accountability, or easier management of subscriptions, etc.

- GeoSynchronization (GSS) - GSS is a mediation service that controls transactional access to one or more WFS's (e.g. to moderate updates in crowdsourcing scenarios). A GSS maintains several event channels, including one for changes applied to the WFS content. Clients can subscribe to the channels (possibly specifying a filter) and be notified by the GSS whenever new entries appear. A GSS may be used to monitor insert/update/delete operations performed on one or more WFS's and send appropriate notifications, implementing the PubSub 1.0 Brokering Publisher Conformance Class. Whenever an event (i.e. a Transaction) occurs on a WFS, the GSS will notify Subscribers of that event. In this way WFS's that do not implement the PubSub 1.0 Standard can participate in an eventing architecture. There are plans to extend GSS to other OGC access services, such as WCS.

[1] <http://www.opengeospatial.org/standards/filter>

[2] <http://www.opengeospatial.org/standards/requests/139>

[3] See also the Proxied Publish/Subscribe use case (access restricted to OGC Members): <https://portal.opengeospatial.org/wiki/PUBSUBswg/PubSubSwgUseCaseBrokeredPubSub>

Chapter 7. PubSub in OGC SensorThings

OGC SensorThings API is a standard for interconnecting the Internet of Things (IoT) devices, data, and applications over the Web. OGC SensorThings API has two main parts, Sensing part which is focusing on heterogeneous IoT sensor systems, and the Tasking part that is focused on IoT actuators and tasking devices.

OGC SensorThings API is following ODATA for managing the sensing resources. Thus, it has a REST-like API and supports HTTP CRUD operations (GET, POST, PATCH, DELETE) as well as ODATA query options (select, expand, filter, orderby, top, skip) for data retrieval.

In addition to supporting HTTP, OGC SensorThings API has the extension for supporting MQTT for creation and real-time retrieval of sensor Observations.

MQ Telemetry Transport, known as MQTT, is an OASIS standard. It is an extremely lightweight publish/subscribe messaging protocol which is specifically designed for resource-constrained IoT devices. The main concepts are topic, and publish and subscribe functions. When a new message is published to a topic, anyone subscribing to that topic will receive a notification including that message.

OGC SensorThings API is adopting MQTT protocol and defines a topic structure for creating and also receiving notifications of sensor Observations. The topic structure for OGC SensorThings API MQTT extension is following its HTTP resource path. It means that the topic used for creating Observation through MQTT is identical to the HTTP URL that is used for POST and creating Observations. Examples of these resource paths are **v1.0/Observations** and **v1.0/Datastreams(id)/Observations**. Similarly, accessing those URLs using HTTP GET would result in retrieving sensor Observations, while subscribing to those as topics means receiving notifications for those Observations in real-time.

Here is the summary of lessons learnt from MQTT adoption is OGC SensorThings API that might be applicable to any RESTful API:

- Each RESTful API has a potential for MQTT binding to receive updates for a resource collection
 - The topic would be the resource GET URL
 - The payload would be the same as content of HTTP GET
 - Whenever there is a new resource, it will be published to the resource GET URL topic
- For any RESTful API, to create a resource, MQTT can be an option just like HTTP POST
 - The topic will be same as the POST topic
 - The payload will be the same as POST payload
 - The service would subscribe to the topics
 - When it receives the payload, it uses the same process as POST for creating the resource

We also did a pilot on a simple rules engine for Institute for the Management of Information Systems (IMIS) pilot project. In that pilot, what we did was to define topics that contain the resource path together with \$filter query options. Then the Observation that are published to that topic would be filtered by the criteria defined in the topic. This way we can define rules that notify

subscribers about certain condition that happened, for example when the CO reading are higher than a threshold and considered dangerous. Example of the topic would be ***v1.0/Datastreams(id)/Observations?\$filter=result gt 100***. This was a pilot implementation and could be a starting point for OGC SensorThings API potential part 3, rules engine.

Chapter 8. W3C Pub Sub Recommendation

Description of the standard and recommendations for OGC PubSub.

Presentation from Stuttgart would be good to draw from.

Chapter 9. Asynchronous Messaging for Aviation from OGC testing

Asynchronous Messaging for Aviation from OGC testing

Chapter 10. WMO OpenWIS use of AMQP

The OpenWIS Association, lead by Meteo France, is examining the use of pubsub services to disseminate Meteorological data as part of the World Meteorological Organization (WMO) Information System (WIS) 2 draft technical specification. The OpenWIS Association conducted an initial WIS 2 study on message protocols and found that several use cases were well positioned to use pubsub. This included but was not limited to the current use of WAN for WMO GTS data dissemination, and use of the WAN for NHMs specific users (aeronautical, B2B and general public).

The WIS 2 study also examined a number of message queue protocols the OpenWIS Program could use including: AMQP, MQTT, STOMP, JMS, Kafka, Redis. However, only 2 were found to use ISO standards (AMQP and MQTT). For a long time, the FAA SWIM aprogram has been very successful in standardizing on the mature JMS API for pub/sub. However, an interoperability limitation in the wire-level protocol for JMS being vendor dependent forced the SWIM program to seek other open source alternatives. SESAR/Eurocontrol is now very much interested in AMQP v1.0 publish/subscribe message protocol as a promising new ISO, IEC, and OASIS international standard, that offers SWIM an open publish-subscribe messaging protocol. Given that, the WIS 2 study focused on specific queues that support AMQP.

The first MQ to be analyzed was RabbitMQ. RabbitMQ supports several message queue protocols including: AMQP 0-9-1, AMQP 1.0 (via a plugin), STOMP, MQTT, HTTP (API). In addition, the study found that RabbitMQ has implemented extensions to AMQP.

In support of our WIS 2 study, the following figure (see figure [Figure 12](#)) shows the proposed AMQP Architecture



Figure 12. OWS Publish/Subscribe FIG1

In our proposed AMQP architecture, all connections are TCP based. There is no UDP. Publish is performed to an exchange, not to a queue. There can be N exchanges. Subscriptions lead to the creation of: 1) a queue (if needed); 2) bindings from exchange to queue. Two basic principles will apply: If there is no binding, there is no routing, and in turn the message is discarded. Also, once messages are consumed, they disappear.

Figure [Figure 13](#) describes one of the proposed AMQP message structure for OpenWIS WIS2



Figure 13. OWS Publish/Subscribe FIG2

Four types of exchanges are being considered for OpenWIS, including Direct, Fanout, Topic, and Header exchanges. Figure [Figure 14](#) describes each of these exchanges

Direct/Fanout/Topic/Header exchanges



Direct exchange: exact routing key match



Fanout exchange: route all



Topic exchange: partial routing key match



Header exchange: routing on header match

Figure 14. OWS Publish/Subscribe FIG3

RabbitMQ tunable settings used during WIS2 study: - Make queues lazy - Store messages on disk (predictable response times) - Make queues persistent - Allows to route messages even if consumer not connected - Tune prefetch - Allows to use all bandwidth even with on high latency WAN - Need to be mindful of big files - Be careful when sending send files bigger than a few Kbs RAM Impact (files loaded to RAM) Transfer times

The proposed message format used during the WIS2 study can be found in figure [Figure 15](#)

```

MANDATORY:
  "pubTime"      - YYYYMMDDTHHMMSS.ss - UTC date/timestamp.
  "baseUrl"      - root of the url to download.
  "relPath"      - relative path can be catenated to <base_url>
  "integrity"    - WMO version of v02 sum field, under development.
  {
    "method" : "md5" | "sha512" | "md5name" | "link" | "remove" | "cod" | "random" ,
    "value"  : "base64 encoded checksum value"
  }
OPTIONAL:
  "size"          - the number of bytes being advertised.
  "blocks"        - if the file being advertised is partitioned, then:
  {
    "method"      : "inplace" | "partitioned" , - is the file already in parts?
    "size"        : "9999", - size of the blocks.
    "count"       : "9999", - number of blocks in the file.
    "remainder"   : "9999", - the size of the last block.
    "number"      : "9999", - which block is this.
  }
  "rename"        - name to write file locally.
  "topic"         - copy of topic from AMQP header (usually omitted)
  "source"        - the originating entity of the message.
  "from_cluster"  - the originating cluster of a message.
  "to_clusters"   - a destination specification.
  "link"          - value of a symbolic link. (if sum starts with L)
  "atime"         - last access time of a file (optional)
  "mtime"         - last modification time of a file (optional)
  "mode"          - permission bits (optional)
  "content"       - for smaller files, the content may be embedded.
  {
    "encoding" : "utf-8" | "base64" ,
    "value"    : "encoded file content"
  }
  For "v03.report" topic messages the following additional
  headers will be present:
  "report" - status report field documented in `Report Messages`_
  "message" - status report message documented in `Report Messages`_
  additional user defined name:value pairs are permitted.

```

Figure 15. OWS Publish/Subscribe FIG4

A single Node pubsub implementation can be found in figure [\[clause7_figure5\]](#).

The results of the study found: 1. When you compare Reference vs Direct message dissemination methods, large message always need to be disseminated by reference 2. For small messages there are questions, including where or not to include messages into payload. It appears to be difficult to set a threshold. This leads to a more complicated protocol, and also leads to less predictable message rate. And finally, it leads to potential issues with memory management 3. The conclusion was to not include message into payload.

The next steps in the study include: 1. Agree pubsub tree. It must be TTAAiiCCCC based. It must be data types based, productor based, or a mix. 2. Need to stabilize the message format 3. Conduct a larger scale test

The key for this is to think about how this could be implemented in the frame of WIS/GTS * Need Pub/Sub for “GISCs” * Need Pub/Sub for “Ncs” * Need SLA / Key Performance Indicators * Need Metadata

Chapter 11. AsyncAPI description

Description of the AsyncAPI spec in general.

The role of AsyncAPI is to provide a language for designing the API.

Chapter 12. Ideas for AsyncAPI in OGC

Application of Async in the context of OGC API work

Chapter 13. Ideas for an OGC Abstract Spec for PubSub

Concepts to be included in an OGC AS for PubSub - draw from the previous sections

Note by Lorenzo.

May look into the chapter in Clause 3 that summarizes the PubSub extension for the generic OWS introduced by OGC 16-137. Such extension is conceived as a simple way to enable the existing request/reply OWS specifications to Publish/Subscribe, by implementing the OGC Publish/Subscribe Interface Standard 1.0.

NOTE

The chapter describes how PubSub 1.0 Core operations, encodings and messages are modeled in terms of the functionalities of the generic OWS. No assumption is made on the capabilities of the target OWS, other than those defined by the OGC Web Services Common Standard. Hence the described extension may apply, for example, to WFS, WCS, and other OWS interfaces.

Annex A: Title

NOTE

Place other Annex material in sequential annexes beginning with "A" and leave final two annexes for the Revision History and Bibliography <<<

Annex B: Revision History

| Date | Release | Editor | Primary clauses modified | Description |
|------------|---------|-----------|--------------------------|-----------------|
| 2016-04-28 | 0.1 | G. Editor | all | initial version |

Annex C: Bibliography

Example Bibliography (Delete this note).

The TC has approved Springer LNCS as the official document citation type.

Springer LNCS is widely used in technical and computer science journals and other publications

NOTE

- For citations in the text please use square brackets and consecutive numbers: [1], [2], [3]

– Actual References:

[n] Journal: Author Surname, A.: Title. Publication Title. Volume number, Issue number, Pages Used (Year Published)

[n] Web: Author Surname, A.: Title, <http://Website-Url>

[1] OGC: OGC Testbed 12 Annex B: Architecture. (2015).

[2] L. Bigagli, M. Rieke. The new OGC Publish/Subscribe Standard - applications in the Sensor Web and the Aviation domain. Open Geospatial Data, Software and Standards, ISSN: 2363-7501, doi: 10.1186/s40965-017-0030-7, vol. 2, issue no. 1, article 18, Springer, Heidelberg, Germany, December 2017 (electronic 3 August 2017).

[3] M. Rieke, L. Bigagli, S. Herle, S. Jirka, A. Kotsev, T. Liebig, C. Malewski, T. Paschke, C. Stasch. Geospatial IoT—The Need for Event-Driven Architectures in Contemporary Spatial Data Infrastructures. ISPRS International Journal of Geo-Information (IJGI), ISSN 2220-9964, Volume 7, Issue 10, article 385, doi: 10.3390/ijgi7100385, MDPI, Basel, Switzerland, 25 September 2018.

[4] OGC. OGC Publish/Subscribe Interface Standard 1.0 — Core, 1st ed.; Open Geospatial Consortium: Wayland, MA, USA, 2016.

[5] OGC. OGC Publish/Subscribe Interface Standard 1.0 - SOAP Protocol Binding Extension, 1st ed.; Open Geospatial Consortium: Wayland, MA, USA, 2016.

[6] Graham, S.; Hull, D.; Murray, B. Web Services Base Notification 1.3; Technical Report; OASIS: Burlington, MA, USA, 2016.

[7] Bigagli, L.; Vretanos, P.P.A.; Lawrence, M.; Papeschi, F.; Martell, R. Testbed-12 PubSub/Catalog Engineering Report; Technical Report OGC 16-137r2; Open Geospatial Consortium: Wayland, MA, USA, 2017.